



Natural Language Processing

Mehmet Can Yavuz, PhD

Adapted from Info 256 - David Bamman, UC Berkeley

History of NLP

- Foundational insights, 1940s/1950s
- Two camps (symbolic/stochastic), 1957-1970
- Four paradigms (stochastic, logic-based, NLU, discourse modeling), 1970-1983
- Empiricism and FSM (1983-1993)
- Field comes together (1994-1999)
- Machine learning (2000–today)

J&M 2008, ch 1

- Neural networks (~2014–today)

Neural networks in NLP

- Language modeling [Mikolov et al. 2010]
- Text classification [Kim 2014; Iyer et al. 2015]
- Syntactic parsing [Chen and Manning 2014, Dyer et al. 2015, Andor et al. 2016]
- CCG super tagging [Lewis and Steedman 2014]
- Machine translation [Cho et al. 2014, Sutskever et al. 2014]
- Dialogue agents [Sordani et al. 2015, Vinyals and Lee 2015, Ji et al. 2016]
- (for overview, see Goldberg 2017, 1.3.1)

Neural networks

- Discrete, high-dimensional representation of inputs (one-hot vectors) -> low-dimensional “distributed” representations.
- Non-linear interactions of input features
- Multiple layers to capture hierarchical structure

Neural network libraries



theano



PYTORCH

py/net

Logistic regression

$$\hat{y} = \frac{1}{1 + \exp\left(-\sum_{i=1}^F x_i \beta_i\right)}$$


	x	β
<i>not</i>	1	-0.5
<i>bad</i>	1	-1.7
<i>movie</i>	0	0.3

SGD

Algorithm 1 Logistic regression gradient descent

- 1: Data: training data $x \in \mathbb{R}^F, y \in \{0, 1\}$
 - 2: $\beta = 0^F$
 - 3: **while** not converged **do**
 - 4: $\beta_{t+1} = \beta_t + \alpha \sum_{i=1}^N (y_i - \hat{p}(x_i)) x_i$
 - 5: **end while**
-

Calculate the derivative of some loss function with respect to parameters we can change, update accordingly to make predictions on training data **a little less wrong next time.**



$$x + a(-2x)$$


$[a = 0.1]$

x	$.1(-2x)$
8.00	-1.6
6.40	-1.28
5.12	-1.02
4.10	-0.82
3.28	-0.66
2.62	-0.52
2.10	-0.42
1.68	-0.34
1.34	-0.27
1.07	-0.21
0.86	-0.17
0.69	-0.14

$$\frac{d}{dx} -x^2 = -2x$$

We can get to maximum value of this function by following the gradient

Logistic regression



$$\hat{y} = \frac{1}{1 + \exp\left(-\sum_{i=1}^F x_i \beta_i\right)}$$


not

bad

movie

x	β
1	-0.5
1	-1.7
0	0.3

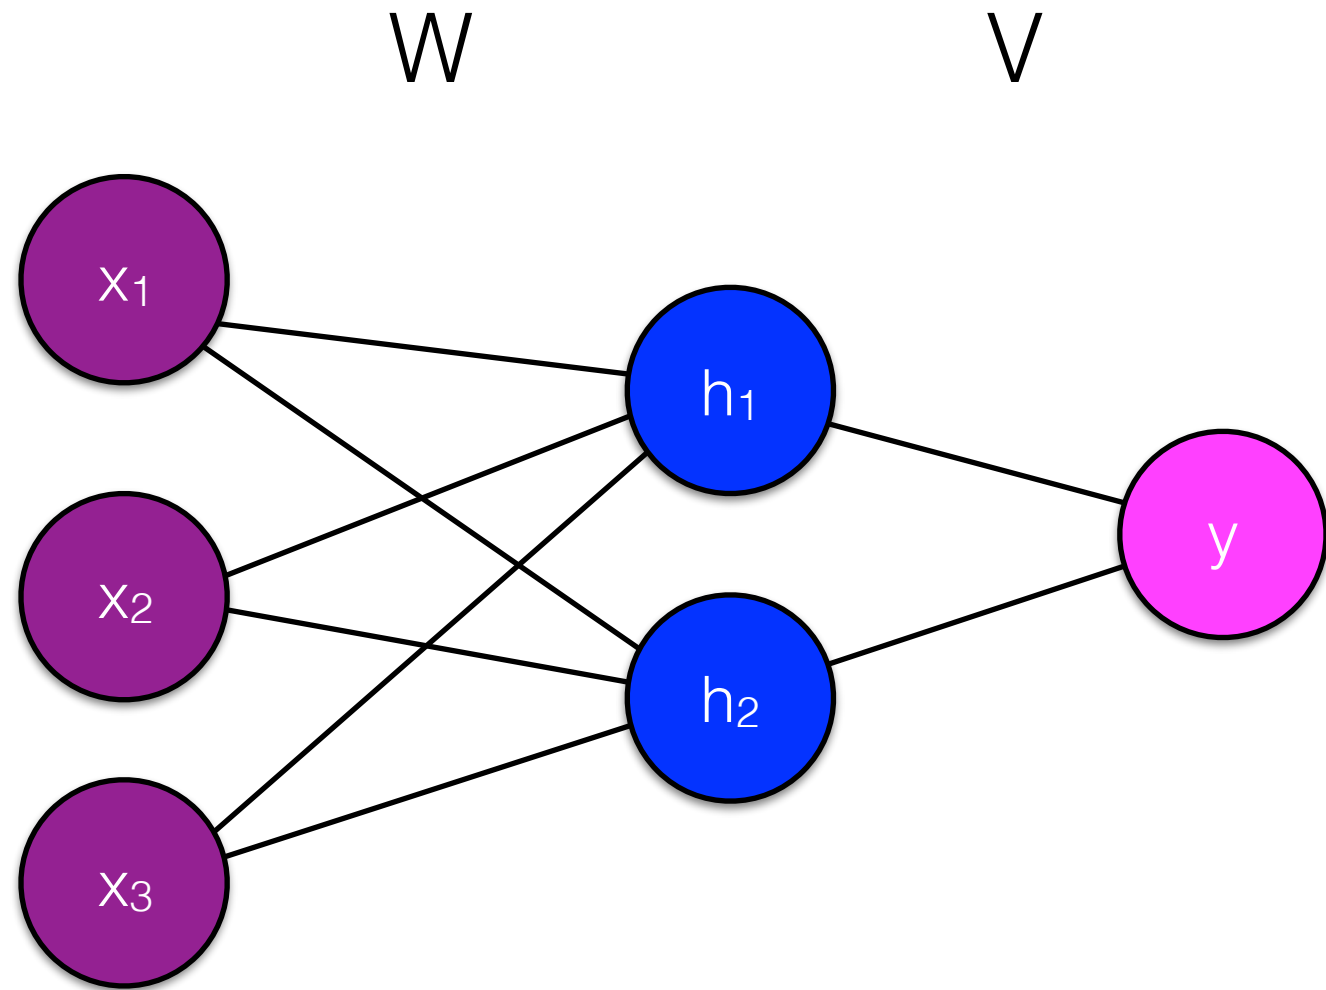
*For simplicity, we're leaving out the bias term, but assume most layers have them as well.



Input

“Hidden”
Layer

Output



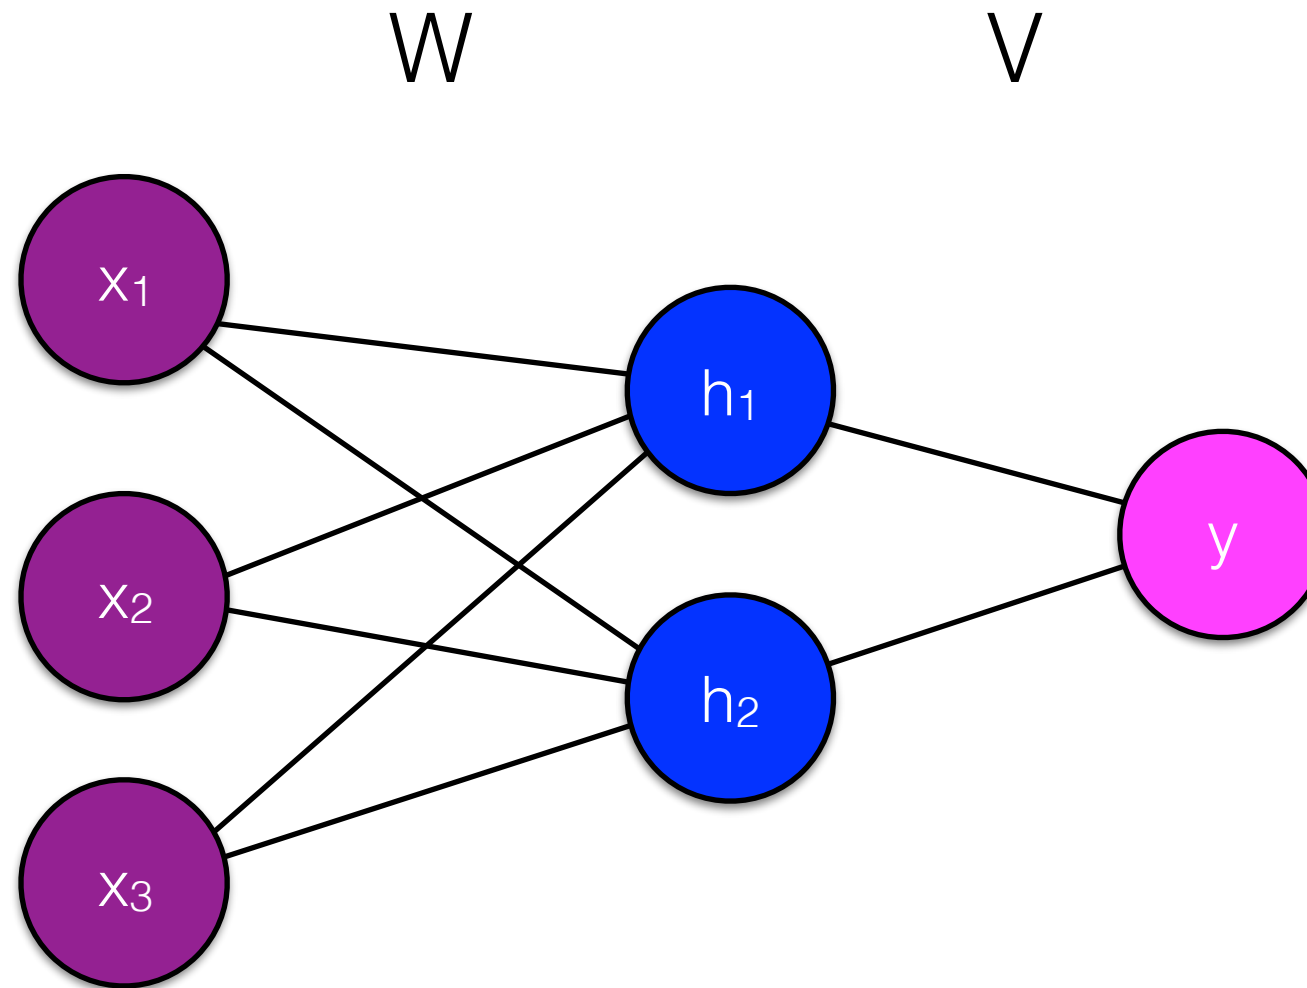
not
bad
movie

x
1
1
0

W	
-0.5	1.3
0.4	0.08
1.7	3.1


V
4.1
-0.9

y
1



$$h_j = f \left(\sum_{i=1}^F x_i W_{i,j} \right)$$


the hidden nodes are
completely determined by the
input and weights




$$h_1 = f \left(\sum_{i=1}^F x_i W_{i,1} \right)$$

Activation functions

$$\sigma(z) = \frac{1}{1 + \exp(-z)}$$



Logistic regression




$$\hat{y} = \frac{1}{1 + \exp \left(- \sum_{i=1}^F x_i \beta_i \right)}$$

$$\hat{y} = \sigma \left(\sum_{i=1}^F x_i \beta_i \right)$$

We can think about logistic regression as a neural network with no hidden layers


Activation functions

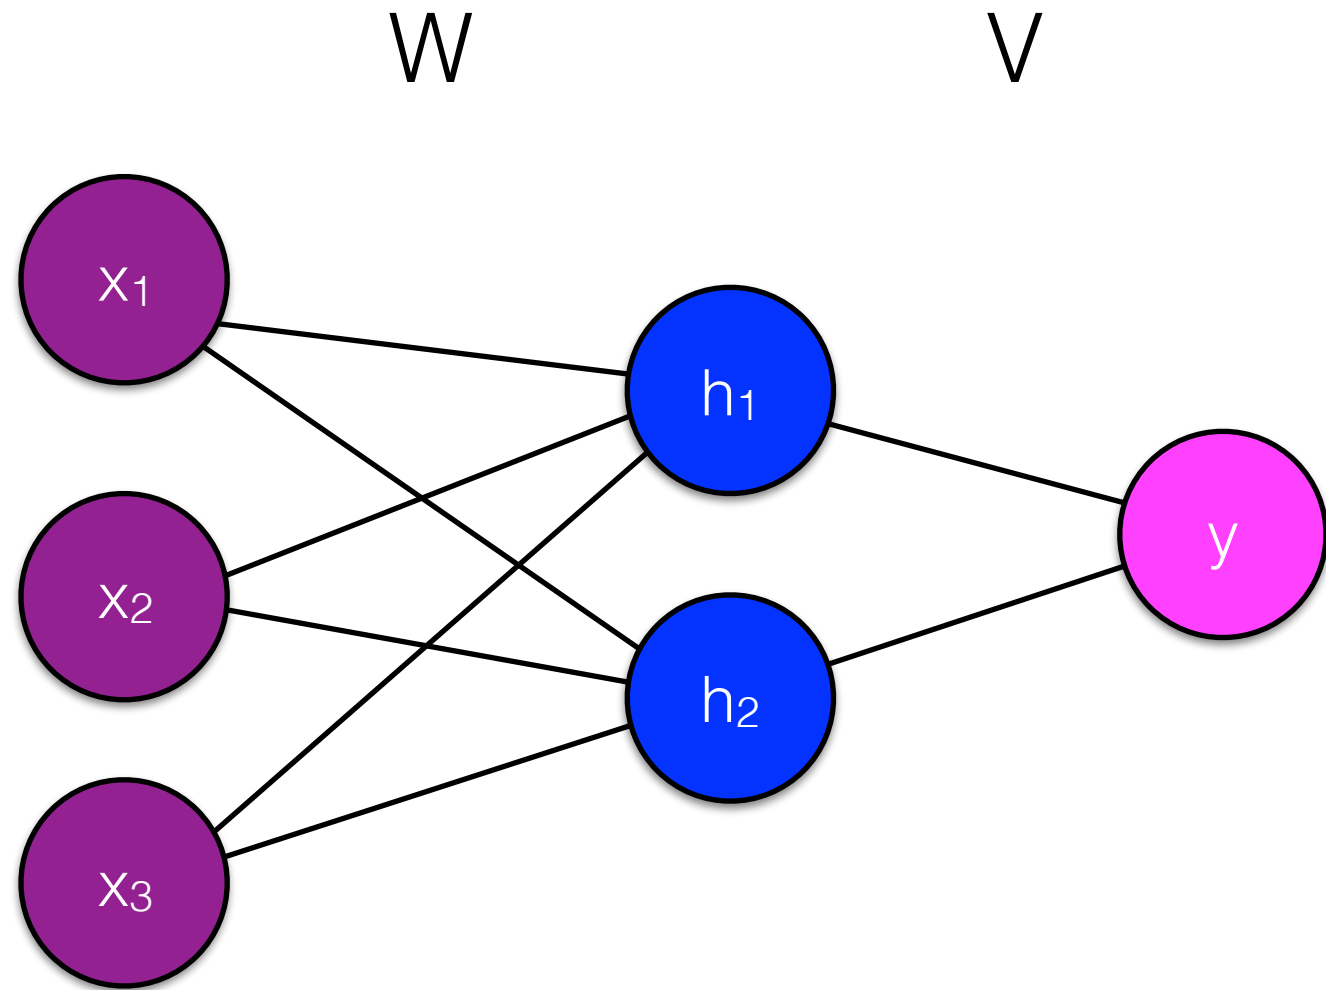
$$\tanh(z) = \frac{\exp(z) - \exp(-z)}{\exp(z) + \exp(-z)}$$



Activation functions

$$\text{rectifier}(z) = \max(0, z)$$






$$h_1 = \sigma \left(\sum_{i=1}^F x_i W_{i,1} \right)$$

$$h_2 = \sigma \left(\sum_{i=1}^F x_i W_{i,2} \right)$$

$$\hat{y} = \sigma [V_1 h_1 + V_2 h_2]$$




$$\hat{y} = \sigma \left[V_1 \left(\sigma \left(\sum_i^F x_i W_{i,1} \right) \right) + V_2 \left(\sigma \left(\sum_i^F x_i W_{i,2} \right) \right) \right]$$

we can express y as a function only of the input x and the weights W and V

$$\hat{y} = \sigma \left[V_1 \underbrace{\left(\sigma \left(\sum_i^F x_i W_{i,1} \right) \right)}_{h_1} + V_2 \underbrace{\left(\sigma \left(\sum_i^F x_i W_{i,2} \right) \right)}_{h_2} \right]$$

This is hairy, but **differentiable**

Backpropagation: Given training samples of $\langle x, y \rangle$ pairs, we can use stochastic gradient descent to find the values of W and V that minimize the loss.



Neural networks are a series of functions chained together

$$xW \rightarrow \sigma(xW) \rightarrow \sigma(xW)V \rightarrow \sigma(\sigma(xW)V)$$

The loss is another function chained on top

$$\log(\sigma(\sigma(xW)V))$$

Chain rule

$$\frac{\partial}{\partial V} \log (\sigma (\sigma (xW) V))$$

Let's take the likelihood for a single training example with label $y = 1$; we want this value to be as high as possible

$$= \frac{\partial \log (\sigma (\sigma (xW) V))}{\partial \sigma (\sigma (xW) V)} \frac{\partial \sigma (\sigma (xW) V)}{\partial \sigma (xW) V} \frac{\partial \sigma (xW) V}{\partial V}$$

$$= \overbrace{\frac{\partial \log (\sigma (hV))}{\partial \sigma (hV)}}^A \overbrace{\frac{\partial \sigma (hV)}{\partial hV}}^B \overbrace{\frac{\partial hV}{\partial V}}^C$$

Chain rule

$$= \overbrace{\frac{\partial \log(\sigma(hV))}{\partial \sigma(hV)}}^A \overbrace{\frac{\partial \sigma(hV)}{\partial hV}}^B \overbrace{\frac{\partial hV}{\partial V}}^C$$

$$= \overbrace{\frac{1}{\sigma(hV)}}^A \times \overbrace{\sigma(hV) \times (1 - \sigma(hV))}^B \times \overbrace{h}^C$$


$$= (1 - \sigma(hV))h$$

$$= (1 - \hat{y})h$$


Neural networks

- Tremendous flexibility on design choices (exchange feature engineering for model engineering)
- Articulate model structure and use the chain rule to derive parameter updates.

Neural network structures




Neural network structures



Multiclass: output 3 values, only one = 1 in training data

Neural network structures



output 3 values, several = 1 in training data


Regularization

- Increasing the number of parameters = increasing the possibility for **overfitting** to training data

Regularization

- L2 regularization: penalize W and V for being too large
- Dropout: when training on a $\langle x, y \rangle$ pair, randomly remove some node and weights.
- Early stopping: Stop backpropagation before the training error is too small.

Deeper networks




Keras

- We'll be using keras to implement several neural architectures over the next few weeks
- Today: Sequential models

Sequential


- Useful for models of limited complexity where the input to every layer is the output of the previous layer.



```
model=Sequential()
```

```
model.add(Dense(2,activation='relu',  
input_shape=(3,)))
```

```
model.add(Dense(1,activation='sigmoid'))
```



```
model=Sequential()
```

```
model.add(Dense(3,activation='relu',  
input_shape=(4,)))
```

```
model.add(Dense(2,activation='relu'))
```

```
model.add(Dense(1,activation='sigmoid'))
```