




Natural Language Processing


Mehmet Can Yavuz, PhD

Adapted from Info 256 - David Bamman, UC Berkeley

Recurrent neural network




Each y is the output of the RNN at that time step; sometimes we use this information (POS tagging, LM); sometimes we only use the output for the final state (s_5)



Each x here is one token in a sequence

Recurrent neural network

- Each time step has two inputs:
 - x_i (the observation at time step i); one-hot vector, feature vector or **word embedding**.
 - s_{i-1} (the output of the previous state); base case: $s_0 = 0$ vector




Recurrent neural network

$$s_i = R(x_i, s_{i-1})$$

R computes the output state as a function of the current input and previous state

$$y_i = O(s_i)$$

O computes the output as a function of the current output state



“Simple” RNN

$g = \tanh$ or relu

$$s_i = R(x_i, s_{i-1}) = g(s_{i-1}W^s + x_iW^x + b)$$


Different weight vectors W transform the previous state and current input before combining

$$W^s \in \mathbb{R}^{H \times H}$$


$$W^x \in \mathbb{R}^{D \times H}$$

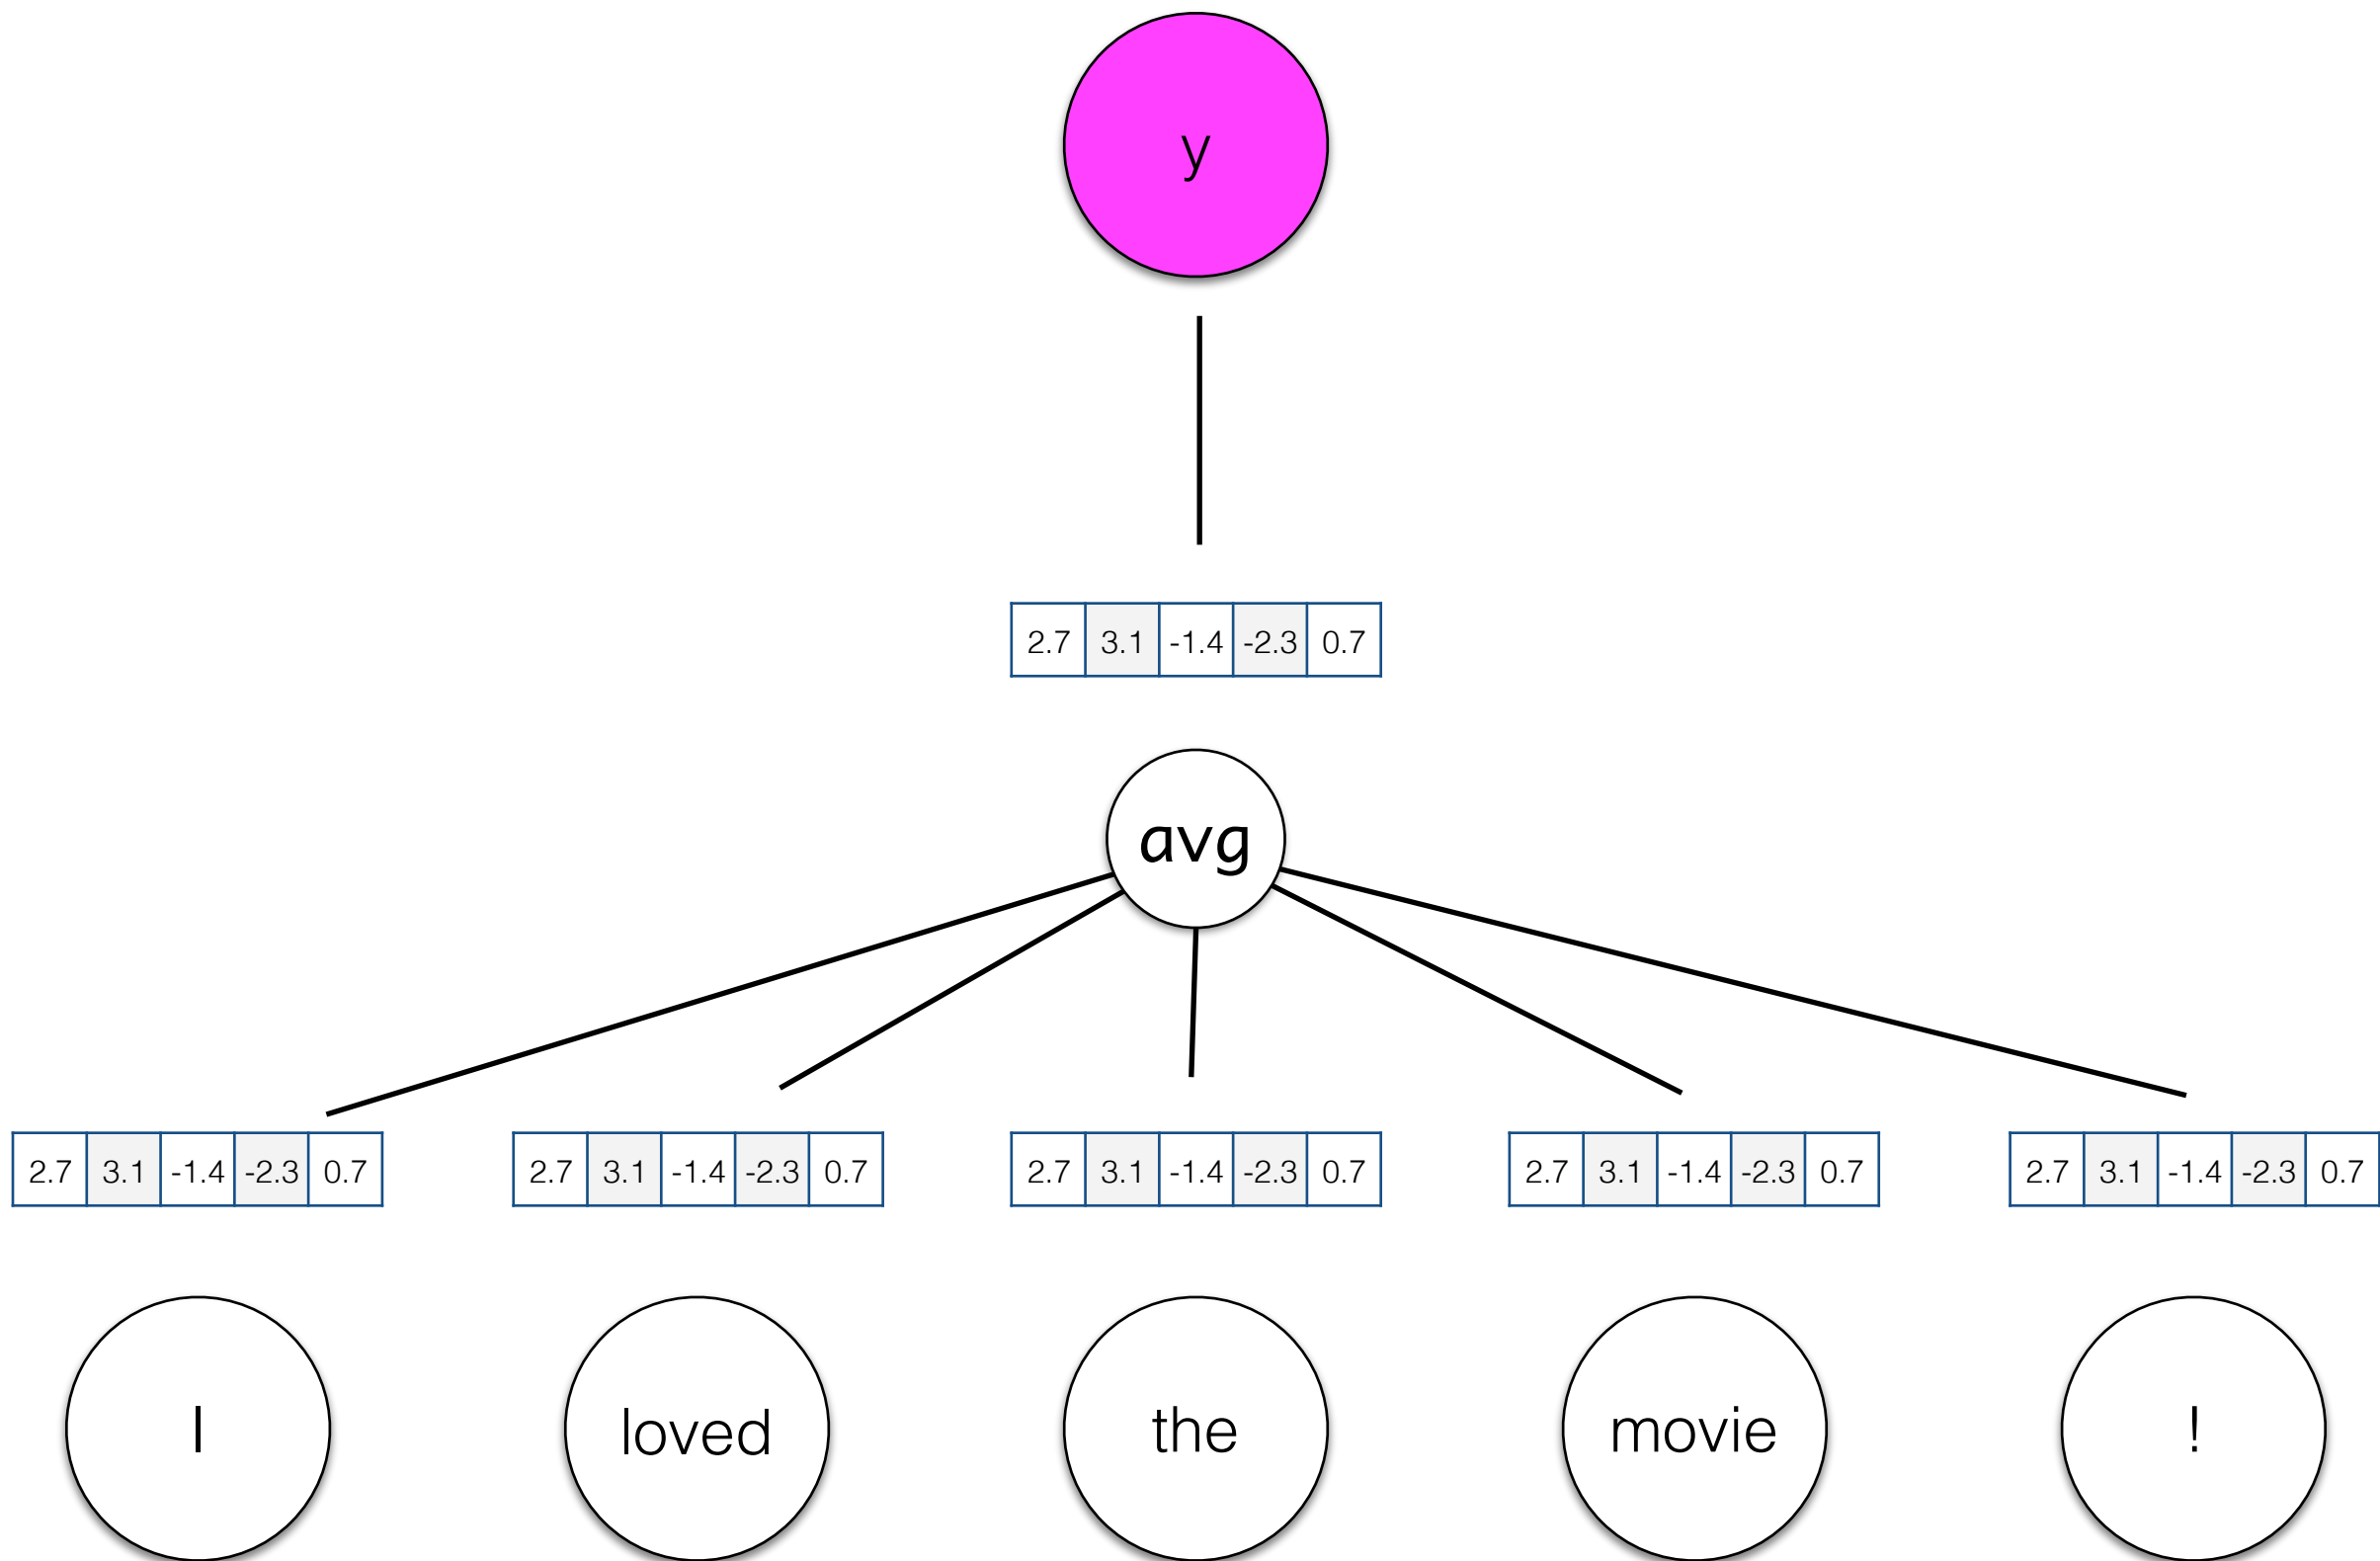
$$b \in \mathbb{R}^H$$

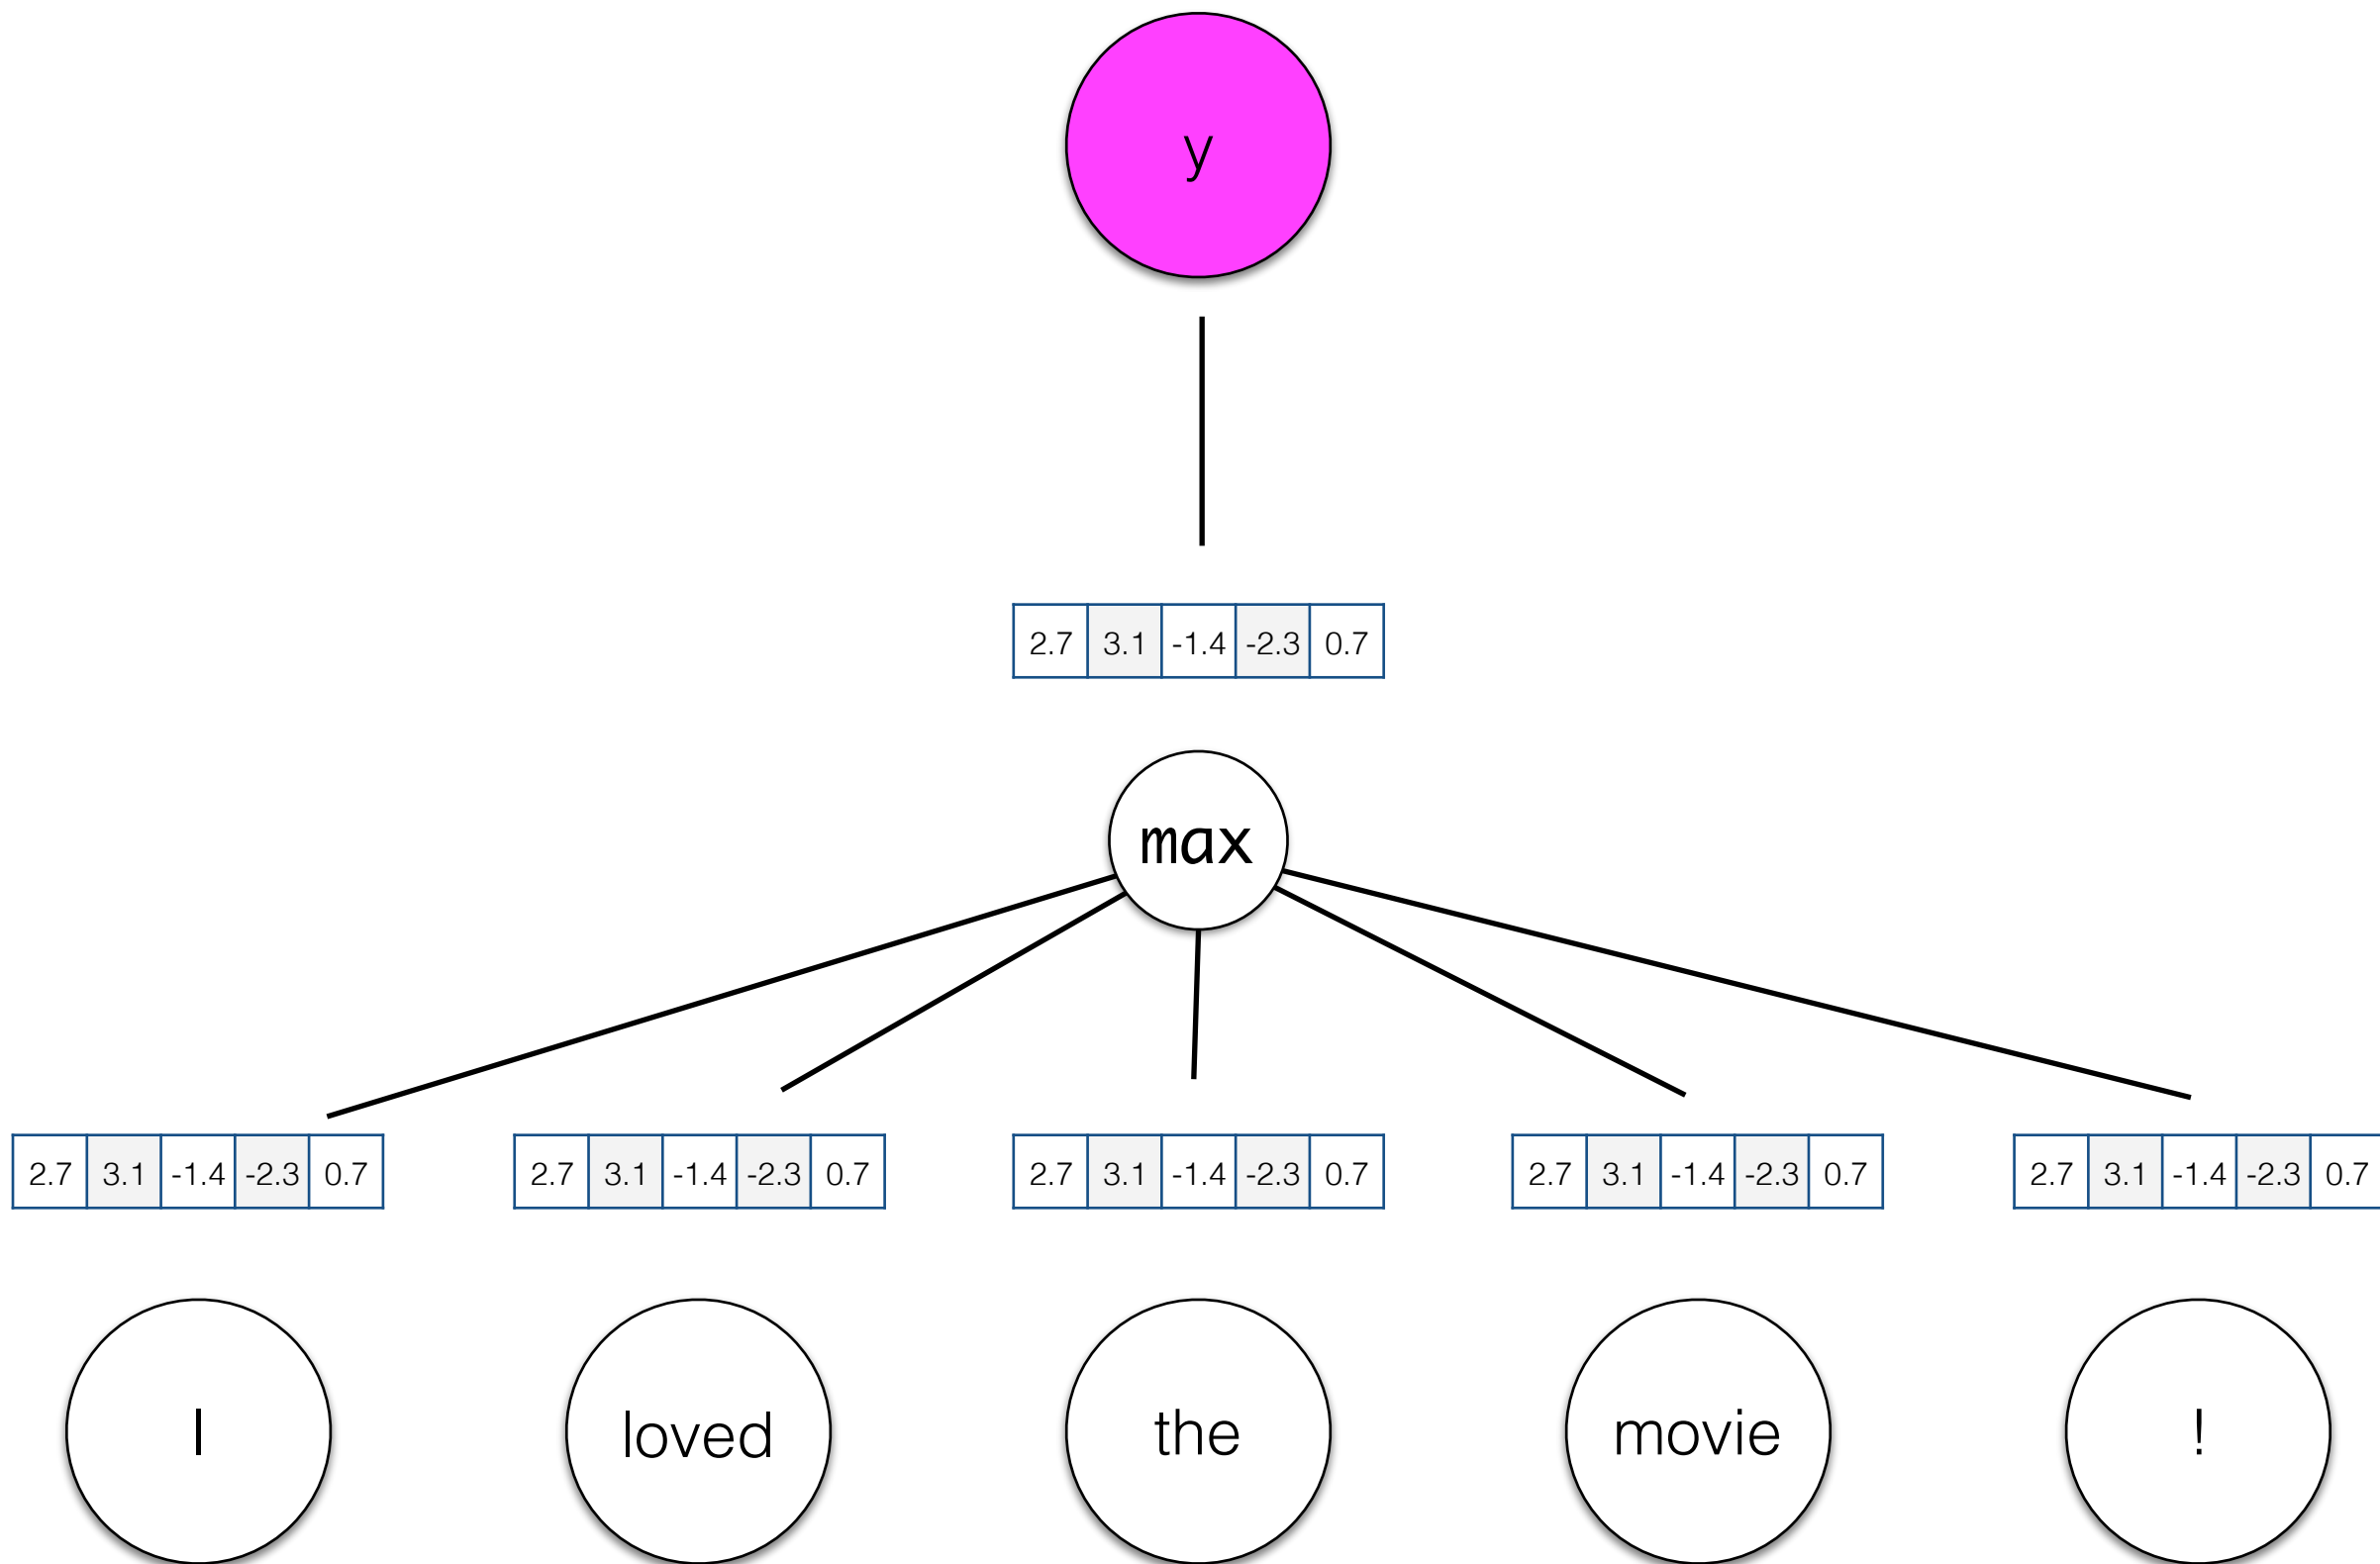
$$y_i = O(s_i) = s_i$$



How do we use RNNs for document classification?

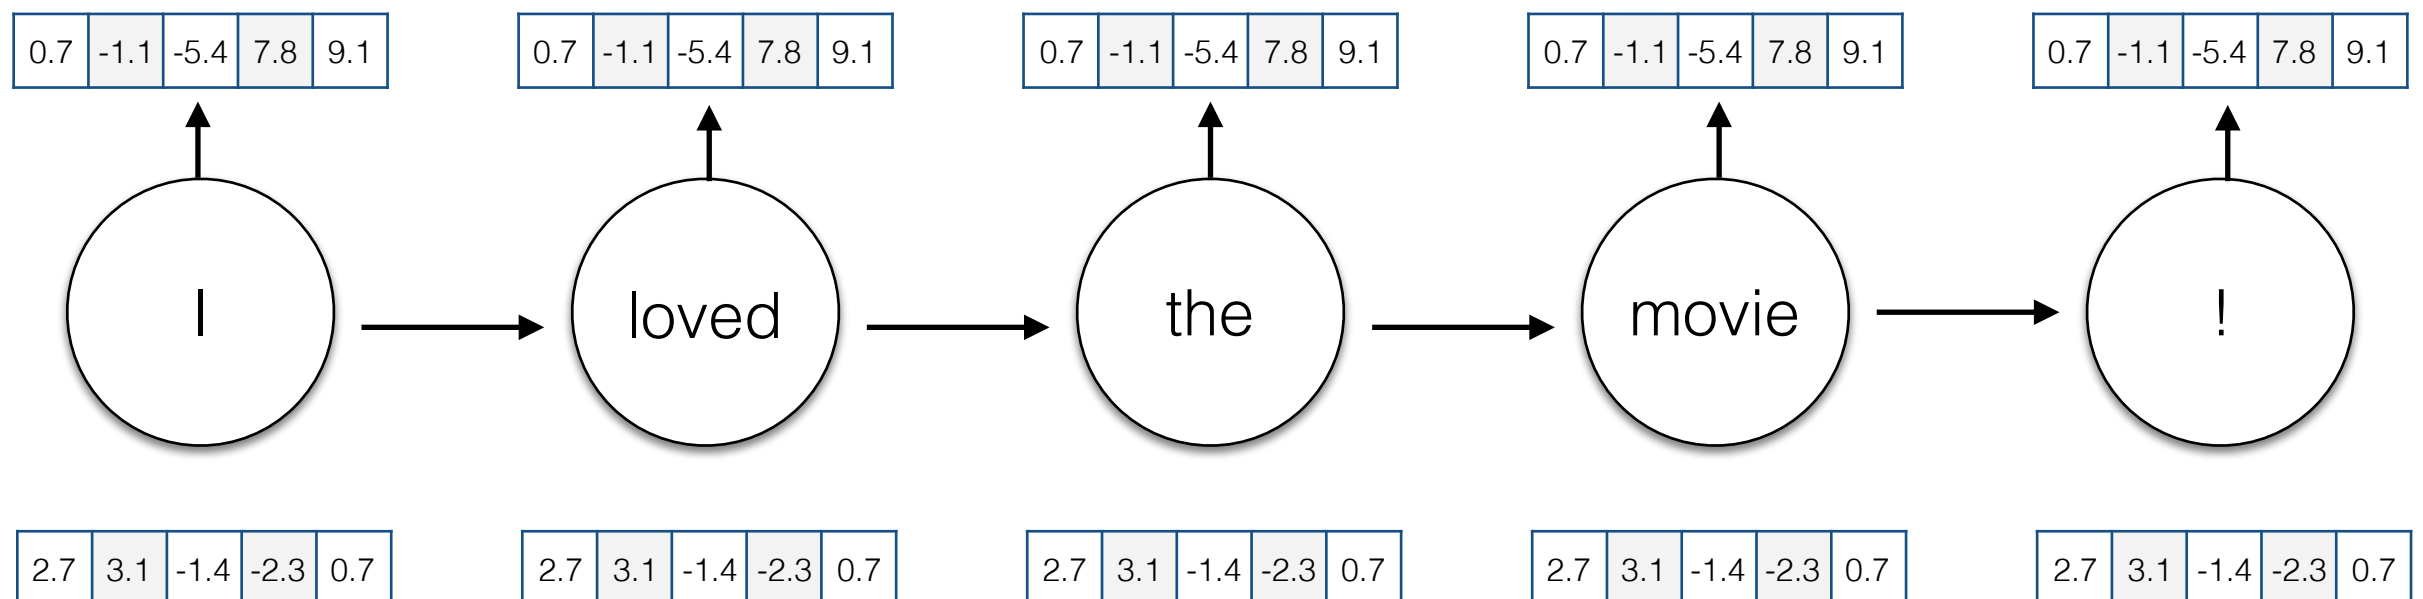




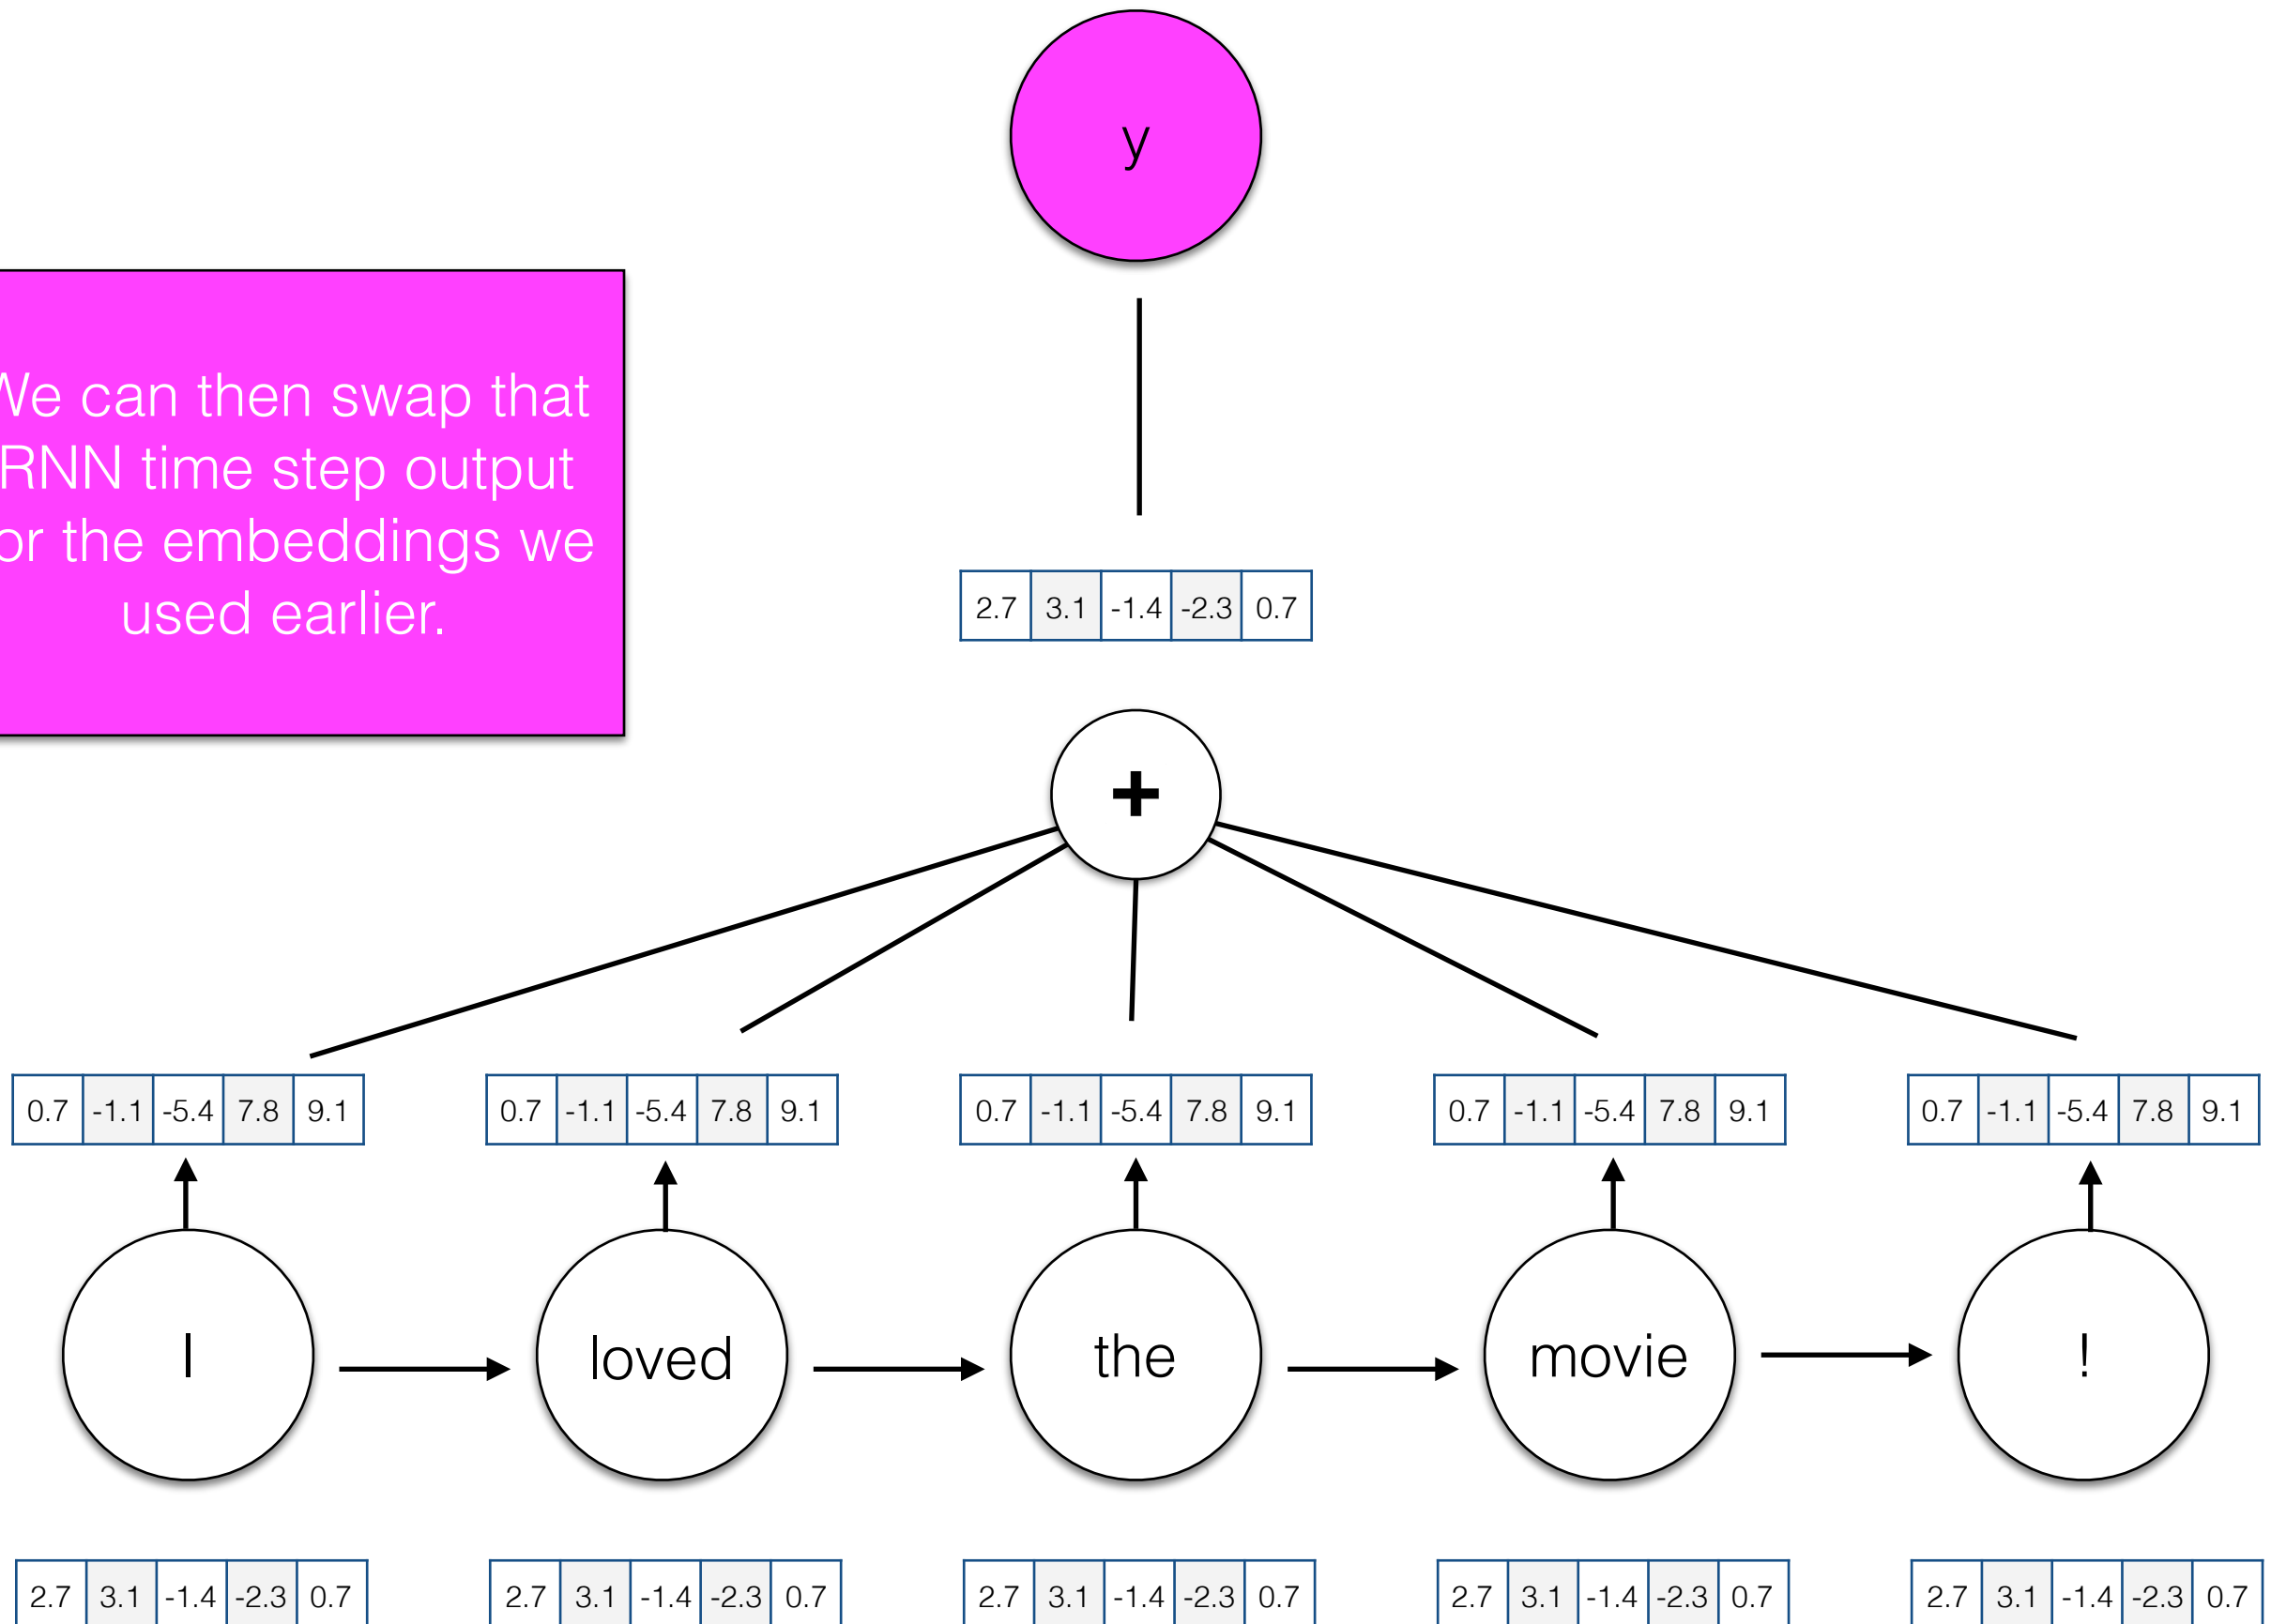


RNN

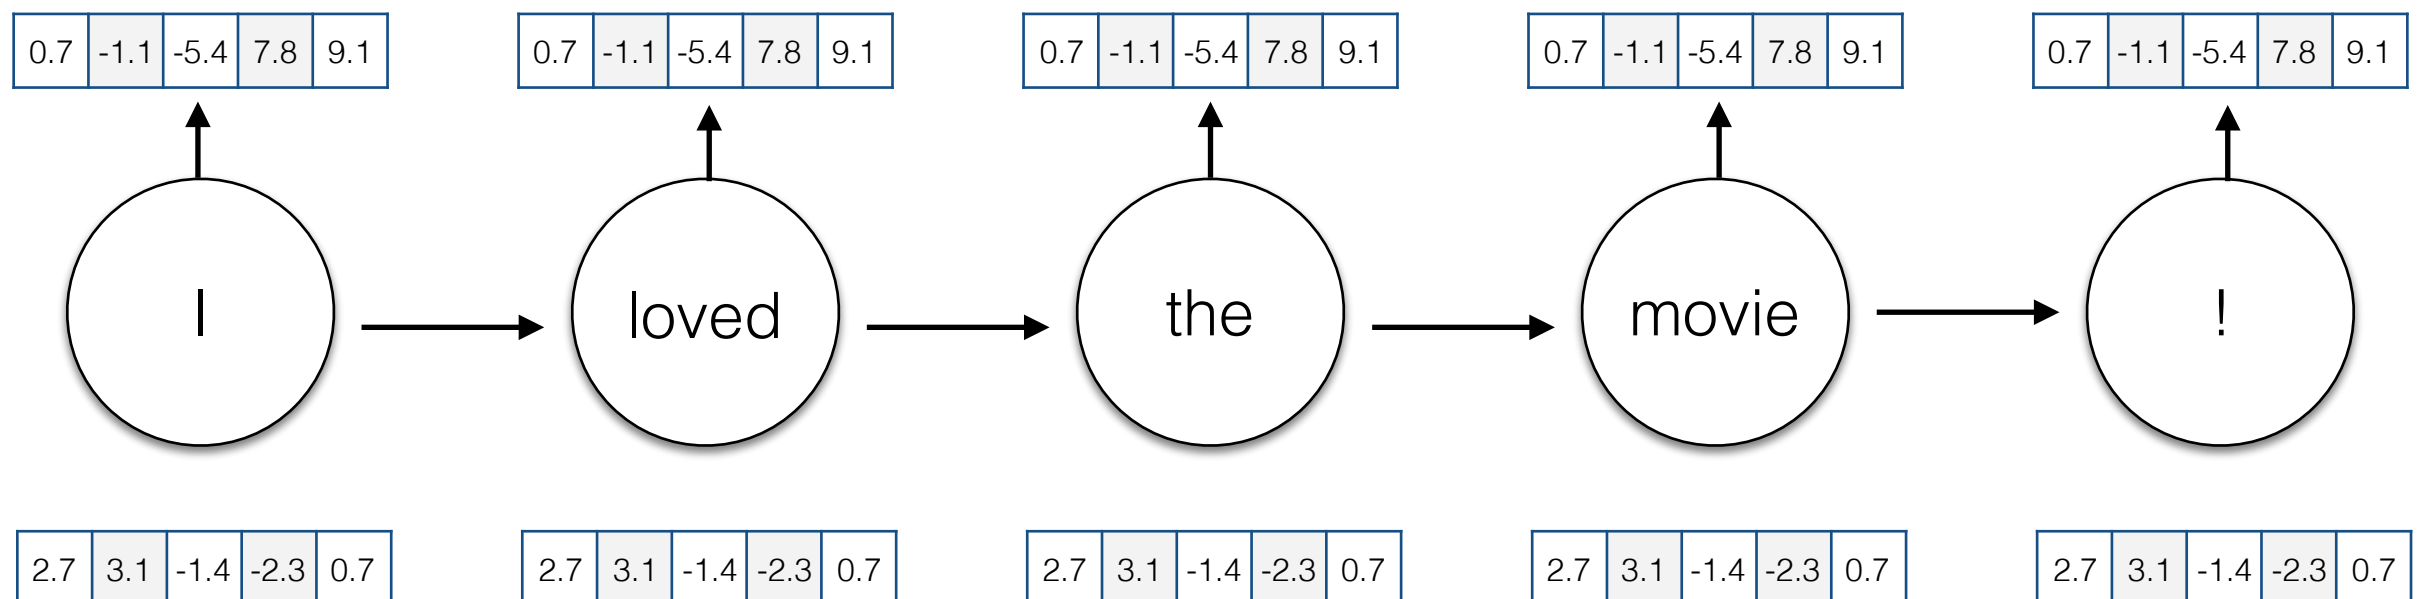
- With an RNN, we can generate a representation of the sequence as **seen through time t**.
- This encodes a representation of meaning specific to the local context a word is used in.




We can then swap that RNN time step output for the embeddings we used earlier.



What about the **future** context?



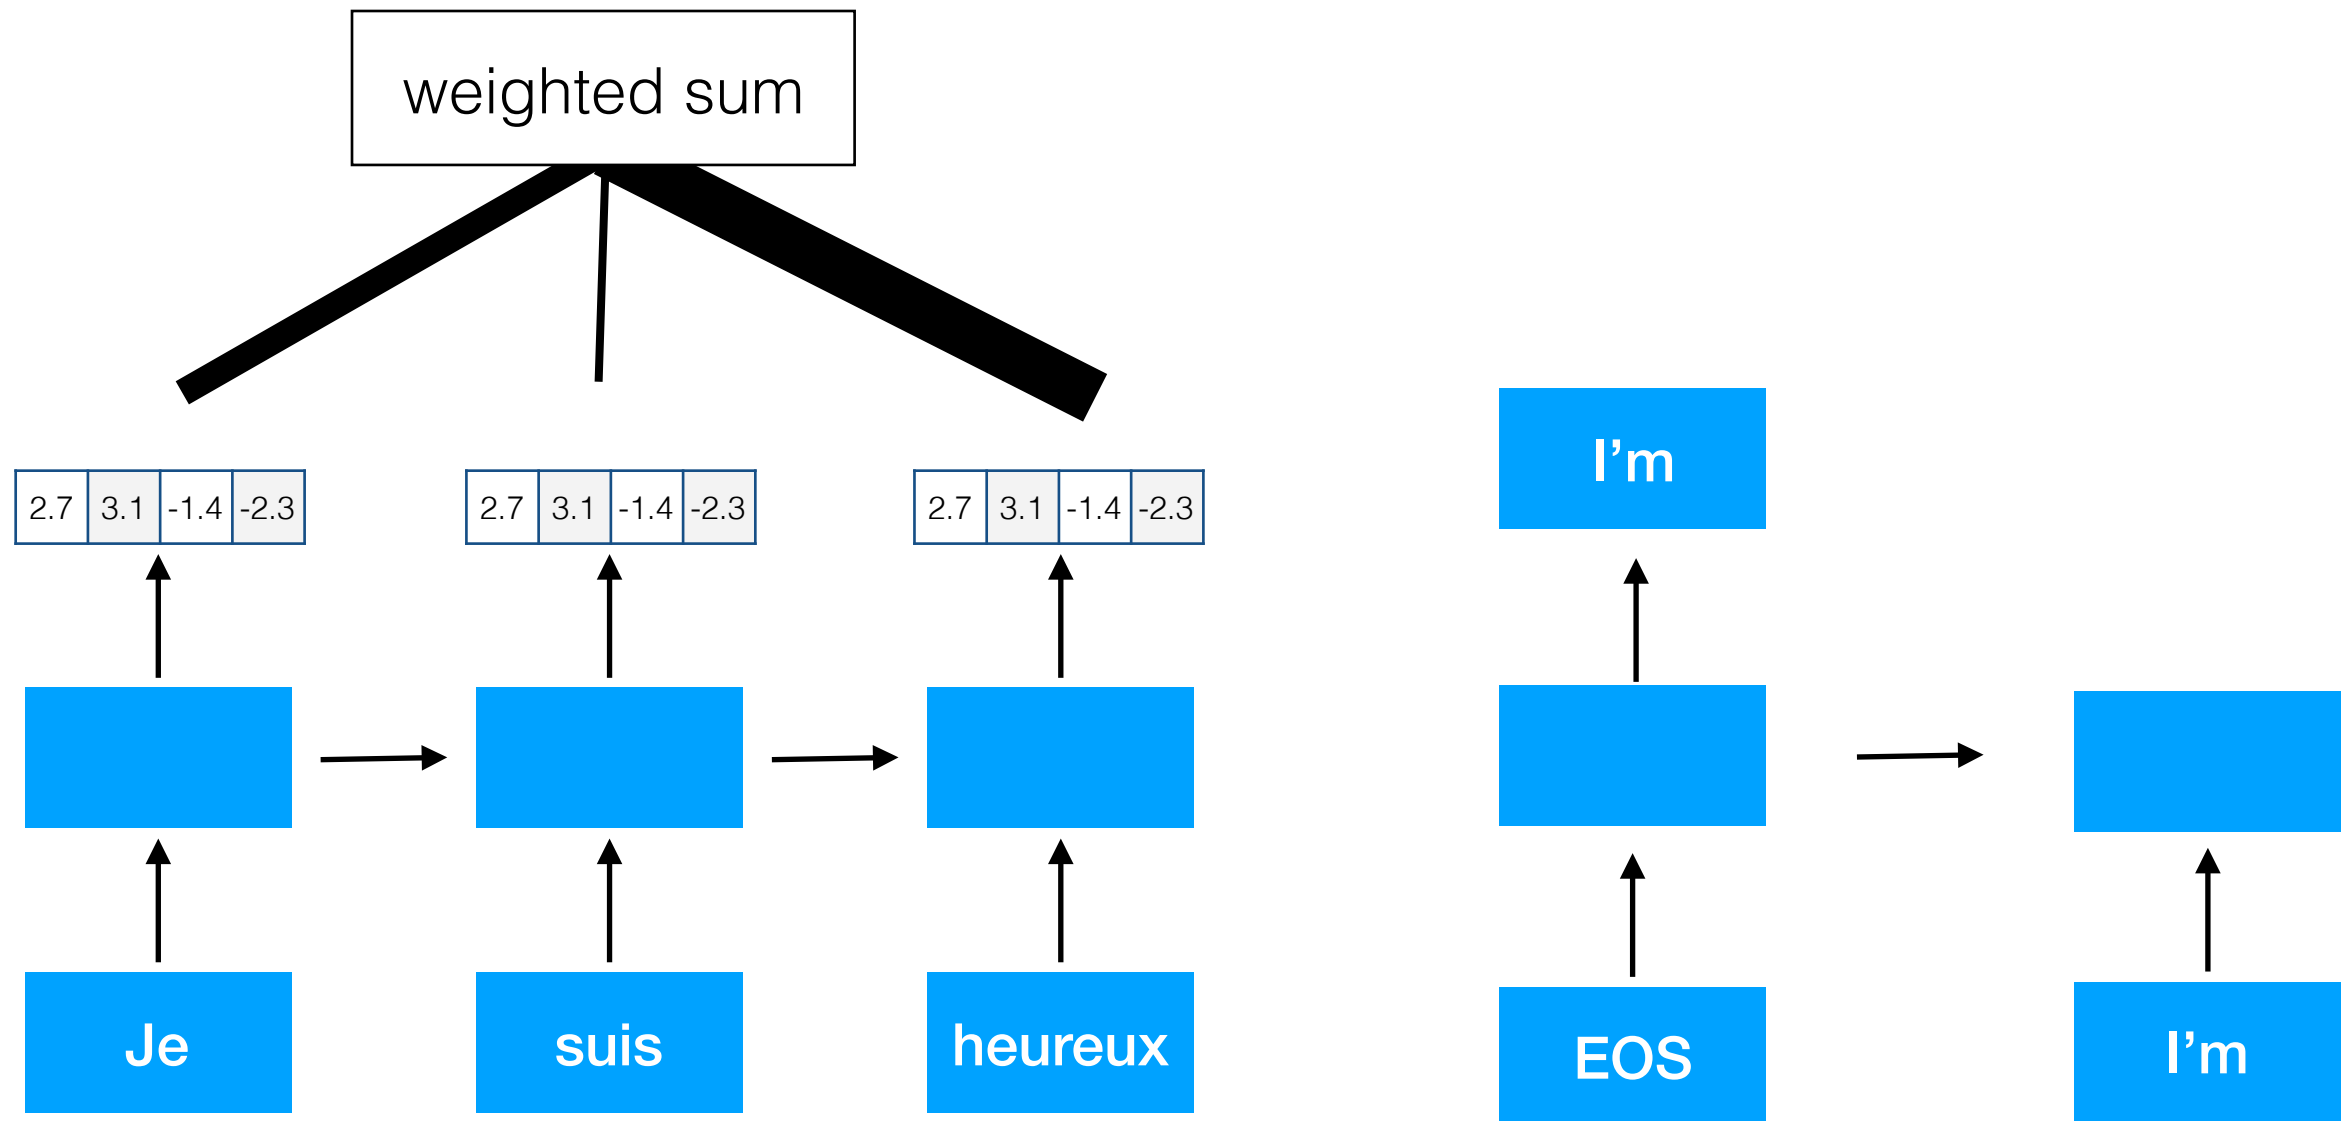
Bidirectional RNN



Attention

- Let's incorporate structure (and parameters) into a network that captures which elements in the input we should be **attending** to (and which we can ignore).

Machine translation



$$v \in \mathcal{R}^H$$

2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----

Define v to be a vector to be learned; think of it as an “important word” vector. The dot product here measures how similar each input vector is to that “important word” vector


2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----

2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----


2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----

2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----


2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----




X_1




X_2



X_3



X_4



X_5

$$v \in \mathcal{R}^H$$

2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----

-3.4

1.7

-0.8

2.4

-1.2

$$r_1 = v^\top x_1$$

|

2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----

$$r_2 = v^\top x_2$$

|

2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----

$$r_3 = v^\top x_3$$

|

2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----

$$r_4 = v^\top x_4$$


|

2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----


$$r_5 = v^\top x_5$$

|


2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----




x_1




x_2



x_3



x_4




x_5

Convert r into a vector of normalized weights that sum to 1.


$$a = \text{softmax}(r)$$

a	0	0.32	0.02	0.64	0.02
r	-3.4	1.7	-0.8	2.4	-1.2


$r_1 = v^\top x_1$	$r_2 = v^\top x_2$	$r_3 = v^\top x_3$	$r_4 = v^\top x_4$	$r_5 = v^\top x_5$																									
<table><tr><td>2.7</td><td>3.1</td><td>-1.4</td><td>-2.3</td><td>0.7</td></tr></table>	2.7	3.1	-1.4	-2.3	0.7	<table><tr><td>2.7</td><td>3.1</td><td>-1.4</td><td>-2.3</td><td>0.7</td></tr></table>	2.7	3.1	-1.4	-2.3	0.7	<table><tr><td>2.7</td><td>3.1</td><td>-1.4</td><td>-2.3</td><td>0.7</td></tr></table>	2.7	3.1	-1.4	-2.3	0.7	<table><tr><td>2.7</td><td>3.1</td><td>-1.4</td><td>-2.3</td><td>0.7</td></tr></table>	2.7	3.1	-1.4	-2.3	0.7	<table><tr><td>2.7</td><td>3.1</td><td>-1.4</td><td>-2.3</td><td>0.7</td></tr></table>	2.7	3.1	-1.4	-2.3	0.7
2.7	3.1	-1.4	-2.3	0.7																									
2.7	3.1	-1.4	-2.3	0.7																									
2.7	3.1	-1.4	-2.3	0.7																									
2.7	3.1	-1.4	-2.3	0.7																									
2.7	3.1	-1.4	-2.3	0.7																									




x_1




x_2




x_3



x_4



x_5



Attention

- For a document with n words and LSTM output e ($= n \times d$ dimensions)
- Dot product between each e_k and attention vector v to yield one r_k for $k = [1, \dots, n]$ ($r = n$ dimensions)
- $a = \text{softmax}(r)$ ($= n$ dimensions)
- Multiply $a * e$ to generate document representation ($= d$ dimensions)


Attention

- Lots of variations on attention:
 - Linear transformation of x into before dotting with v


Pass the LSTM output through a mini neural network to generate r

0.1	0.5	0.1	0.2	0.1
-----	-----	-----	-----	-----

$$a = \text{softmax}(r)$$



Apply the attention weights to the original LSTM outputs



2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----

$$x_1a_1 + x_2a_2 + x_3a_3 + x_4a_4 + x_5a_5$$

weighted sum






2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----

2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----

2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----

2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----

2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----



Attention

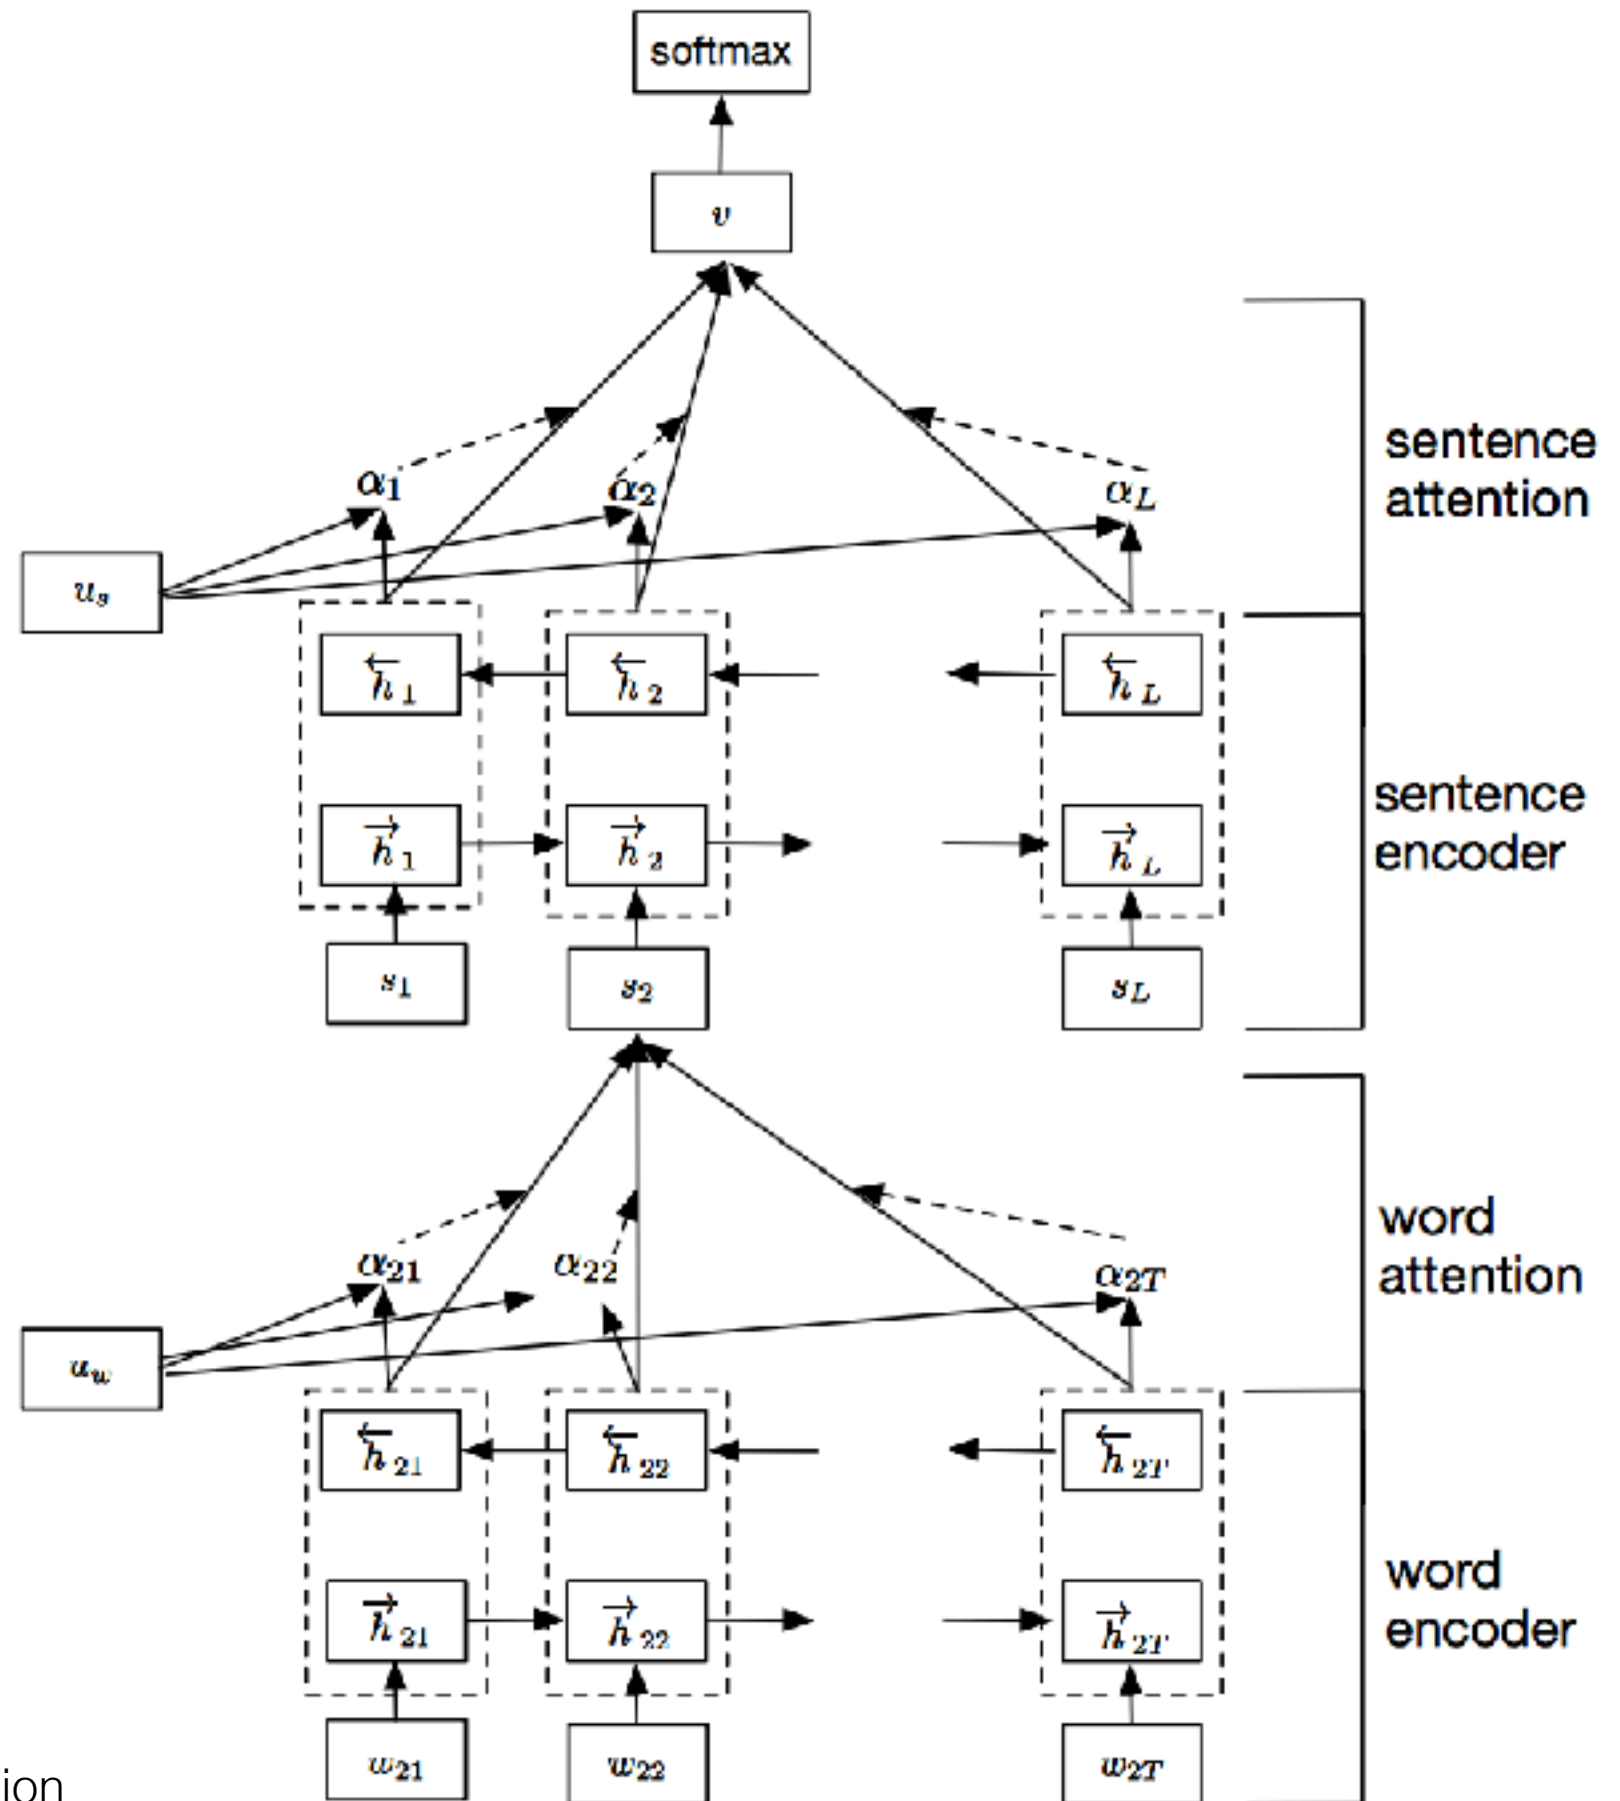
- Lots of variations on attention:
 - Linear transformation of x into before dotting with v
 - Non-linearities after each operation.
 - “Multi-head attention”: multiple v vectors to capture different phenomena that can be attended to in the input.
 - Hierarchical attention (sentence representation with attention over words + document representation with attention over sentences).

attention over sentences

bidirectional GRU over
sentence representations

attention over words

bidirectional GRU over
word representations



Attention

- Attention gives us a normalized weight for every token in a sequence that tells us how important that word was for the prediction
- This can be useful for visualization

